ARTICLE

PRISM: making glioma progression scores inspectable on MRI after treatment

Sri Datta Ganesh Bandreddi[a] and Jessica Zhang[b]

[a]Department of Biomedical Engineering, Carnegie Mellon University, Pittsburgh, PA, USA; [b]Department of Mechanical Engineering, Carnegie Mellon University, Pittsburgh, PA, USA

ARTICLE HISTORY

Compiled November 13, 2026

ABSTRACT

After treatment for aggressive glioma, changes on serial MRI must be read alongside the treatment course. A progression score can rank surveillance scans, but cannot show whether it reflects MRI measurements or treatment context. We developed PRISM, a local system that makes progression scores inspectable by combining tumour measurements and treatment context in a checked report. On MU-Glioma-Post, the public development cohort, the combined model separated scans with recorded progression from other scans (AUC 0.82). Applied without refitting to external cohorts, prespecified models reached AUCs of 0.77 at LUMIERE (combined) and 0.80 at UCSF-ALPTDG (images from one scan). Audits showed the scores followed treatment timing at MU and enhancing volume at UCSF. PRISM shows each score with the measurements, treatment context and provenance behind it. After one revision of the prompt, all 406 reports met predefined automatic checks against their structured records. Future work will evaluate the reports with clinicians.

KEYWORDS

glioma surveillance, external validation, conformal prediction

1. Introduction

After surgery and chemoradiotherapy for glioblastoma and other aggressive gliomas, serial MRI is read against earlier scans and the course of treatment. One early challenge is pseudoprogression. Treatment can produce new or increasing enhancement that resembles tumour growth. Pseudoprogression is concentrated within 12 weeks of completing chemoradiotherapy (Wen et al. 2023) and is more common in tumours with MGMT promoter methylation (Taal et al. 2008; Brandes et al. 2008).

Models intended to predict progression assign a score to each surveillance scan. On its own, a score gives an assessment without showing the information that informed it. A score does not show whether the available information agrees, whether relevant information is missing, or how the model has performed in comparable clinical situations.

CONTACT Sri Datta Ganesh Bandreddi. Email: sbandred@andrew.cmu.edu; ORCID: 0009-0004-3617-6734

PRISM is a system that makes progression scores inspectable. For each scan, it keeps image measurements, treatment context, and separate model scores together in a record that is checked before it is written as a report. The name comes from this separation. A prism splits light, and PRISM keeps image measurements and treatment context apart. Earlier systems have generated brain tumour reports from structured features extracted from segmentation (Heras Rivera et al. 2026; Valerio et al. 2025), and deterministic checks have compared generated clinical text with source records (Zeba et al. 2026). PRISM applies these ideas to glioma surveillance after treatment. Its language model writes the completed record as text, and an automatic checker compares the text with that record.

We developed PRISM on the public MU-Glioma-Post cohort and validated it, without refitting, in the LUMIERE and UCSF-ALPTDG cohorts. Testing it showed why a score should be inspectable, because the same kind of score rested on different information in different cohorts. This paper makes four contributions. The first two are the methods we built, and the last two are what testing them showed.

(1) This paper introduces PRISM, a system that builds one record for each follow up scan. The record holds tumour volumes and their change measured from segmentation masks, three separately fitted scores (a clinical-history model, an imaging model and a combined model), and three statements of confidence (an 80% conformal set, whether the clinical-history and imaging models agree, and how the models performed in comparable situations), each value with its source. We defined the rules for agreement and for comparable situations after seeing the first results (Sections 2.1–2.3, validated in Sections 4.1–4.3).
(2) This paper presents a reporting workflow that runs entirely on local hardware. A language model writes the report from the record alone, citing for each sentence the record fields it uses, and a deterministic checker tests every number, flag and confidence statement in the report against the record (Section 2.4, validated in Section 4.4).
(3) This paper shows that what a progression score rests on depends on how a cohort documented progression. At MU the scores followed treatment timing, and at UCSF enhancing volume alone ranked scans as well as any fitted model, which an AUC alone cannot reveal (Section 4.5).
(4) This paper tests PRISM in two external cohorts with models fitted once on MU and, for its planned tests, rules fixed in advance. We test whether the scores still rank scans, whether the confidence statements keep their stated coverage, and whether every report passes the checker (Sections 3.2 and 4.2–4.4).

Section 2 describes PRISM, from its inputs and measurements (Section 2.1) and its three models (Section 2.2) to its confidence statements (Section 2.3), its record and report (Section 2.4), and what it costs to run (Section 2.5). Section 3 describes the three cohorts, the task, how each cohort was evaluated, which tests were fixed in advance, and the audits of what carries a score. Section 4 validates the inputs, the three models, the confidence statements and the reports on MU and, with the models locked, in the two external cohorts. It ends with the audits of what information carries the scores. Section 5 discusses what the results mean for anyone who builds or evaluates such scores, and what remains untested, above all whether the reports help clinicians.

2. PRISM

In five stages (Figure 1A), PRISM turns one follow up scan and the patient's clinical record into a structured record and a report. PRISM measures the tumour, scores the scan with three models, attaches three confidence statements, and writes every value into one record with its source. A local language model then writes the report from that record, and a checker compares the report with the record. Table 1 lists every score used in the paper.

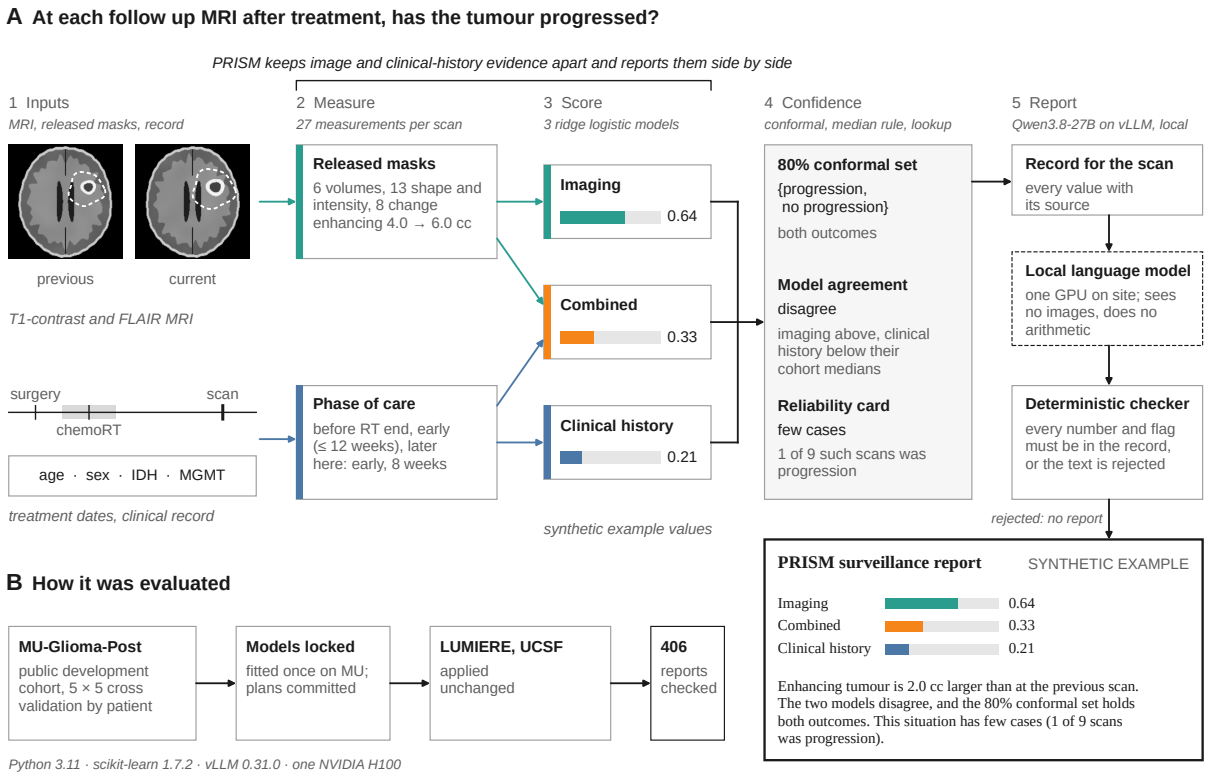


Figure 1. PRISM and how it was evaluated. (A) At each follow up MRI, PRISM keeps image and treatment evidence apart. The masks released with the cohort give the volumes of the previous and current scans and 27 measurements of the tumour and its change, which feed the imaging model (three volumes and their change at LUMIERE). The treatment dates and clinical record give the position in treatment, which feeds the clinical-history model. A third model on both inputs gives the combined score, and all three are ridge logistic regressions (values from the synthetic example of Figure 4). Three confidence statements are attached, a split conformal set at the 80% level, the agreement of the two models and a lookup of results in the same situation. Every value enters one record with its source, a local language model (Qwen3.8-27B served by vLLM on one GPU) writes the report from the record without seeing images, and a deterministic checker tests the report against the record. (B) Models were developed on MU by repeated cross validation, then fitted once and applied unchanged at LUMIERE and UCSF under rules fixed before any score was compared with an external label.

2.1. Inputs and measurements

PRISM takes three inputs for a scan. The first is the segmentation of the current scan and of earlier scans, which gives the volumes of enhancing tumour, of FLAIR hyperintensity, and of necrosis and the resection cavity. We used the images and masks as released with each cohort (Mahmoud et al. 2025; Suter et al. 2022; Fields et al. 2024), already registered to a common space by each provider,

Table 1. The scores built from measurements and records. Every score listed is a ridge logistic regression except enhancing volume alone, which is not fitted. The others are fitted within the training data, with unknown IDH or MGMT status as the reference category and missing measurements imputed with the training median and an indicator. The networks tuned on the task are described in Section 3.4. Status indicates whether a model was planned before the analysis or not fixed in advance (†).

Score	Inputs	Dates	Status
Clinical data alone	Age, sex, IDH and MGMT status	No	Planned reference (LUMIERE †)
Clinical-history model	Clinical inputs; days since diagnosis and surgery; radiotherapy start, end and dose; chemotherapy; visit index and previous scans	Yes	Planned
Reduced clinical-history model	Clinical inputs; days since surgery; flag for the first weeks after radiotherapy	Partly	Planned at UCSF-ALPTDG; † at LUMIERE
Imaging model	27 measurements of the segmented tumour: volumes, shape and intensity, change since the previous scan, volumes relative to the patient's smallest	No	† on MU and in its three volume form at LUMIERE; planned at UCSF-ALPTDG and in the system tests
Images from one scan	The 19 measurements of the current scan	No	† on MU; planned at UCSF-ALPTDG
Combined model	Inputs of both models; externally, the reduced clinical-history inputs, volumes and their change, and the interval since a previous scan	Yes; partly externally	† on MU (full form); planned and locked (reduced form) at LUMIERE and UCSF-ALPTDG
Planned primary model	Clinical-history inputs, image measurements and three schedule inputs (interval since the previous scan, whether there was one, smallest volume since the end of radiotherapy)	Yes	Planned primary model on MU; not part of PRISM
Enhancing volume alone	Log enhancing tumour volume, no fitted model	No	† reference

so every volume can be traced to a public file. PRISM does no registration, resampling or segmentation of its own. From the masks and images PRISM computes 27 measurements. Nineteen describe the current scan. They are the log volumes of the six compartments (enhancing tumour, necrosis, FLAIR hyperintensity, resection cavity, tumour core and whole lesion) and 13 measures of shape and of T1 contrast and FLAIR intensity, the intensities divided by the median of normal appearing brain more than 10 mm from the lesion and cavity (Supplementary Table S47). The other eight describe change. They are the log ratios of the volumes to those of the previous scan and the enhancing and lesion volumes relative to the smallest the patient has had up to that scan (Table 1). The second input is the treatment context as of the scan day. It records the days since diagnosis and surgery and the start, end and dose of radiotherapy and chemotherapy. From it PRISM places the scan in one of four phases of care, before the end of radiotherapy, early after it (up to 12 weeks after its end), later, or unknown when the dates are missing. The third input is the patient's age, sex, and IDH mutation and MGMT promoter methylation status, two tumour markers recorded at diagnosis. Every input is taken as of the scan day, so no information from a later visit enters a score, except IDH and MGMT status at 8 grade 4 scans taken before the diagnosis date (Supplementary Section S4).

2.2. Three models

PRISM scores each scan with three models. The clinical-history model uses the treatment context and the patient characteristics, the imaging model uses the 27 measurements and no dates or intervals, and the combined model uses both sets of inputs. Each is a logistic regression with an L2 (ridge)

penalty. We chose it because its coefficients can be read and because deep networks fine tuned on MU added nothing to the clinical-history model (Table 4). Missing values are imputed with the median of the training scans and flagged with an indicator, unknown IDH or MGMT status is coded as the reference category, and the inputs are standardised. The penalty is chosen from six values of the inverse strength C (0.01, 0.03, 0.1, 0.3, 1 and 3) by threefold cross validation grouped by patient within the training data, minimising the log loss. Each model outputs a score between 0 and 1 that ranks scans by how closely they resemble progression scans in the training data.

Where a cohort lacks treatment dates, PRISM uses a reduced clinical-history model with the inputs that cohort has (Table 1). The flag for the first weeks after radiotherapy is taken from the rater's note at LUMIERE and approximated at UCSF as 42 to 125 days after the radiotherapy start, assuming a six week course. LUMIERE's masks have three compartments and no separate cavity, so for LUMIERE we fitted a separate imaging model on MU with only those three volumes and their change since the previous scan[†]. UCSF holds two scans per patient, so there the change measurements use the first scan as the only earlier one. Because the imaging model needs only scans and masks, PRISM still gives an imaging score in a cohort with no treatment dates.

2.3. Confidence statements

PRISM attaches three confidence statements to each score and keeps them separate, so that each can be checked on its own.

The first is a conformal set at the 80% level (Angelopoulos and Bates 2023), which names the outcomes the combined score cannot rule out. We refitted the combined model on 70% of the MU training patients and used the other 30% for calibration. Write $\hat{p}(1) = \hat{p}$ for a scan's score and $\hat{p}(0) = 1 - \hat{p}$. For each calibration scan i with outcome y_i, the nonconformity score is $s_i = 1 - \hat{p}_i(y_i)$, the probability the model did not give to what happened. Because a patient contributes several scans that are not independent, each scan is weighted by one over the number of that patient's calibration scans, so that every patient carries the same weight. With $\alpha = 0.2$ and n calibration patients, the threshold $\hat{q}$ is the smallest s_i whose cumulative weight, with the scores sorted, reaches $(1-\alpha)(n+1)$. The added one counts the new patient with weight one. If no score reaches it, $\hat{q}$ is infinite. A new scan with score $\hat{p}$ receives the set

$$C(\hat{p}) = \{\, y \in \{0, 1\} : 1 - \hat{p}(y) \le \hat{q} \,\}. \tag{1}$$

The set holds one outcome when the score is clear, both when it is not, and in principle none when the threshold is very small. We computed the set from the model refitted on 70% of the patients instead of the displayed combined score. This costs some discrimination (at UCSF the refitted model's AUC was 0.08 below that of the displayed score). The record shows the combined model's set. We computed sets in the same way for the clinical-history and imaging models to test them. If a new patient is exchangeable with the calibration patients, meaning drawn from the same population in no particular order, the set contains the true outcome of about 80% of their scans on average, but not necessarily within each subgroup (Romano et al. 2020). On MU the calibration split was made inside each training fold. Externally we used one model and one threshold from MU, unchanged.

The second statement records whether the clinical-history and imaging models agree. The two models agree when both scores fall on the same side of their medians over the cohort's evaluation

scans. A score equal to the median counts as below. This rule uses no outcome labels. Because a cohort median needs a batch of scans, we later also tested medians fixed on MU, which can be applied to a single scan† (a dagger marks an analysis not fixed in advance, Section 3.3).

The third statement is a reliability card† (Mitchell et al. 2019), a lookup of how the models performed on scans in the same situation. A situation is the phase of care crossed with the state of the enhancing tumour. That state is absent when the scan has less than 0.1 cc, unknown without a usable segmentation, without comparison when there is no previous scan, increased when it grew by at least 1 cc and 40% from a measurable previous volume or by at least 1 cc of new enhancement from none, and not increased otherwise. For each cohort the card counts the scans with and without progression in that situation among the cohort's held out results. With at least 10 of each it shows the AUC of all three models with 95% confidence intervals (tested), with fewer it shows the counts only (few cases), and with none it says the situation was not tested. The card's counts include the scan being described, so it is optimistic. PRISM also flags a scan whose IDH or MGMT status is unknown or whose inputs lie outside the 1st to 99th percentile of the training data.

2.4. Record and report

Every value PRISM computes for a scan enters one structured record under a fixed key, with its source. A local language model (Qwen3.8-27B, run with greedy decoding) then writes the report from that record alone, so no patient data are sent to an outside service. The prompt is assembled from the fields present in the record, so the model is never asked about information the record lacks, and decoding is constrained to a JSON schema in which every sentence carries its section and the keys of the record fields it uses. The model sees no images and does no arithmetic, and every number it writes must match a field it cites, up to the rounding, sign and percentage forms the record already implies.

A deterministic checker then accepts or rejects the whole report. The checker rejects a report whose output cannot be parsed, or in which a sentence cites a field that does not exist, states a number not found in a field it cites, or writes a number as a word. The checker also rejects a report in which a raised flag is not mentioned, an input is wrongly called out of range, or the conformal set or the agreement statement is missing or misstated, or the reliability tier is missing where the evidence is thin. It also rejects a report that leaves out any of the three scores or the enhancing volume, that quotes a performance figure absent from the reliability card, or that offers a diagnosis, a response category, advice, certainty words such as "rules out" or a verbal probability such as "likely". A rejected report is not repaired. No rule ties a number to the label it is written next to or checks the direction of a change, and the checker cannot judge emphasis or what a clinician would infer (Supplementary Section S19).

No public cohort releases the clinicians' reports for these scans, so we compared the report's content with the follow up template of the Brain Tumor Reporting and Data System (BT-RADS) (Weinberg et al. 2018), which clinicians found less ambiguous than free prose (Gore et al. 2019; Zhang et al. 2020). Like the template, the report gives the position in treatment and the change in enhancing tumour and FLAIR abnormality since the previous scan, here as measured volumes. The report leaves out the template's category and management advice, which PRISM does not give, and steroid and bevacizumab use, which the cohorts do not record (Supplementary Section S24).

2.5. Implementation and computational cost

PRISM is written in Python 3.11 with scikit-learn 1.7.2 (Pedregosa et al. 2011), and its language model is served by vLLM 0.31.0 (Kwon et al. 2023) in bfloat16 precision on one NVIDIA H100 GPU with 80 GB of memory. The model weights take 50.2 GiB. On one core of an Intel Xeon Platinum 8481C processor, computing the 27 measurements from a scan's masks and images took about 2 s (median of 20 MU scans). Fitting one of the three models with its penalty search on synthetic data of MU's shape (181 scans) took under 0.1 s. Scoring took about 2 µs per scan when scans were scored together, and computing the conformal threshold took under 1 ms. Writing reports took 1.08 s per report when generated in a batch, after about two minutes to load the model. A scan therefore costs about 2 s of processor time and 1 s of GPU time once segmented. These costs leave out segmentation, because we used the masks released with each cohort. A deployment would add the time of an automatic segmentation model. Every stage is one command of a single package with fixed random seeds, and each run records the hashes of its inputs, code and outputs (Supplementary Section S16).

3. Evaluation design

3.1. Cohorts and the task

Grade 4 glioma is the primary population in all three cohorts (Table 2). The task is to decide whether each scan is the one at which the cohort documented progression. Predicting progression within 120 days of a scan is a secondary task. On MU the intended use covers scans from the start of radiotherapy onwards. MU results refer to the intended use unless stated otherwise.

3.2. How each cohort was evaluated

Each cohort played one role (Table 3). MU was the development cohort. On it we evaluated every model by five repeats of fivefold cross validation grouped by patient, so that no patient contributed to both fitting and testing. We report the mean pooled held out AUC. We used repeated splits because single splits are unstable in small cohorts (An et al. 2021; Varoquaux 2018) (Supplementary Section S8). We then fitted each model once on all 181 grade 4 detection scans of MU, from every phase of care, locked it, and applied it unchanged to LUMIERE and UCSF under plans committed before any score was compared with an external label. The plans named all 229 MU detection scans of gliomas of every grade. Fitting on those instead of the grade 4 scans changed the LUMIERE AUC by less than 0.01. As a reference, we also fitted each model on UCSF itself by cross validation.

Discrimination is measured by the AUC, the probability that a scan at which progression was documented scores higher than a scan without progression. Every AUC carries a 95% confidence interval (CI) from 1,000 bootstrap samples of patients, and on MU these intervals resample held out predictions without the variance from refitting. Conformal sets are judged by coverage, the share of scans whose set contains the true outcome, against a planned threshold of 78% at the 80% level. The agreement statement is judged by whether the combined score ranks scans less well where the two models disagree, and by whether setting those scans aside helps more than removing as many scans matched for phase of care. Reports are judged by the share that pass the checker, by whether every report in a situation with few cases or none says so, and by whether reports follow a deliberately

Table 2. The three cohorts. What each cohort records, and the scans in the detection task. Grade 4 glioma is the primary population.

	MU-Glioma-Post	LUMIERE	UCSF-ALPTDG
Access	Public	Public	Data use agreement
Scans per patient	Up to 6 scans with a label over 120 days (median 2)	Up to 10 scans with a label over 120 days (median 2)	One consecutive pair
Treatment dates	Surgery, radiotherapy start and end, dose, chemotherapy	Days from surgery only; rater flag for scans within three months of radiotherapy	Days from surgery; radiotherapy start only
Documented progression (the label)	Documented date in the clinical record	Expert response rating of each scan	Progression date supplied with the dataset
Segmentation masks	Automatic masks released with the dataset	DeepBraTumIA (three labels)	Masks provided with the dataset
Patients[a]	126	61	140
Age, median	62.5 (IQR 53–70)	62 (IQR 53–70)	57.5
Male	66 (52%)	32 (52%)	81 (58%)
IDH wildtype / mutant / unknown	112 / 8 / 6	38 / 1 / 22	81 / 13 / 46
MGMT methylated / unmethylated / unknown	45 / 62 / 19	26 / 28 / 7	71 / 43 / 26
Detection task[b]: scans with / without progression (patients)	62 / 84 (80) in the intended use; 64 / 117 (92) over all scans	73 / 105 (79)	29 / 40 (69), second scan only
Phase of care in the detection set	Progression: 56 of 62 later after radiotherapy; no progression: 51 of 84 early after radiotherapy	Not recorded	65 of 69 later after radiotherapy[c]

[a] Patient characteristics describe the cohorts of the secondary task over 120 days at MU-Glioma-Post (grade 4) and LUMIERE, and all grade 4 patients at UCSF-ALPTDG. Age is at diagnosis at MU-Glioma-Post and at surgery at LUMIERE.

[b] Positives are the scan at which progression was documented, and negatives are defined per cohort in Table 3. The LUMIERE detection set has more patients than the cohort over 120 days because it includes patients who contribute only a progression scan. At MU-Glioma-Post, the 146 scans of the intended use are the 181 grade 4 detection scans less the 35 taken before radiotherapy started, and 42 grade 4 progression scans of patients with no scan carrying a label over 120 days are excluded. Reports were evaluated on the scans of patients not used to develop the prompt, which at UCSF-ALPTDG include grade 4 second scans outside the detection set (Section 4.4).

[c] UCSF-ALPTDG gives only the first radiotherapy start, so the phase assumes a six week course. Two scans were early after radiotherapy and two had no radiotherapy date.

Table 3. How each cohort was evaluated. Definitions of the label and the procedure in each cohort. Scan counts are in Table 2.

	MU-Glioma-Post	LUMIERE	UCSF-ALPTDG
Role	Development	External test	External test
How models were applied	Five repeats of fivefold cross validation grouped by patient; then fitted once on all grade 4 detection scans and locked	Locked MU models, applied once	Locked MU models, applied once; models refitted on UCSF only as a reference
Scan with progression	First scan on or after the documented progression date	First scan rated as progressive disease	Second scan of the pair, with progression dated within 30 days and closer to it than to the first scan
Scan without progression	Followed by at least 120 days without progression	Earlier scan rated stable or responding	Second scan followed by at least 120 days without progression
Scores evaluated	All scores in Table 1	Clinical data alone and the reduced clinical-history, imaging and combined models	Clinical data alone and the reduced clinical-history, imaging model, images from one scan and combined model
Plan committed before the test	Detection plan (2b4c927); system plan (e7cba80)	Detection plan (2b4c927); system plan (e7cba80)	UCSF plan (0e7c9fe); system plan (e7cba80)

changed record. To test the checker itself, we introduced one error at a time into correct reports written for synthetic records and counted how many it rejected, under a plan written before the run.

3.3. What was tested, and the rules fixed in advance

Each planned test answered its own question, so we made no correction for multiple tests. Table 4 lists every part of the record, the test that checks it, the rule fixed in advance where there was one, and the outcome. Results of analyses not fixed in advance carry a dagger (†). The disagreement flag and its two tests were planned, but we wrote down the median rule that defines disagreement after seeing the results. Every disagreement result therefore carries a dagger.

3.4. Audits of what carries a score

Each label records how a cohort documented progression, so a score can partly reproduce the way progression was called rather than what the scan shows. On MU, 48 of the 64 scans with progression were taken on the day progression was documented. We used four audits to see what information carries a score. The first is the AUC within phases of care, which pools only pairs of a scan with progression and a scan without progression from the same phase, so that treatment timing cannot separate them. The second is enhancing tumour volume alone, with no fitted model, as a reference for how far a label follows enhancement. The third builds the imaging model in three steps, from the current scan's measurements to change from earlier scans and then three schedule inputs, and compares the AUC across and within phases at each step. The fourth tests whether deep networks detect more in the images than the measurements capture. We fine tuned SwinUNETR segmentation networks to score scans directly and compared them with the same networks trained on permuted labels (Supplementary Section S22). Only the fourth audit was fixed in advance.

Table 4. What was tested, and the outcome. Each part of PRISM's record, the test that checks it, the rule committed to the repository before the test where there was one, and the outcome. AUC differences with 95% CIs. † not fixed in advance.

Part of the record and test	Rule fixed in advance	Outcome	Section
Inputs			
Volumes, agreement of a new MU segmentation with the released masks	New masks raise detection by at least 0.05, CI above 0	**No**. Median enhancing tumour Dice 0.80, no gain, no expert reading	4.1
Phase of care	None	From treatment dates on MU, the rater's note at LUMIERE, approximated at UCSF. Not checked	2.1
The three models			
The planned primary model adds to the clinical-history model on MU	Planned primary model ≥ 0.75 (lower bound ≥ 0.65); gain ≥ 0.03, CI above 0	**No**. AUC 0.837 (0.763–0.904), gain +0.015 (−0.032 to +0.063)	4.2
Deep networks tuned on MU add	Gain ≥ 0.03, CI above 0, and control with permuted labels covers 0.5	**No**. No network gained. Control 0.483 (0.436–0.527)	2.2
Combined score at LUMIERE	AUC ≥ 0.65	**Yes**	4.2
Imaging model at UCSF	AUC ≥ 0.65, lower bound > 0.5	**Yes**	4.2
Locked imaging model adds to clinical data alone, both locked, at UCSF	Gain, CI above 0	**Yes**. +0.248 (+0.096 to +0.392)	4.2
Deep networks learn at UCSF	Above the control with permuted labels; control covers 0.5	**Yes**. +0.201 (+0.094 to +0.306) over a control of 0.541 (0.480–0.603)	4.5
Confidence statements			
Conformal set, coverage in each cohort	≥ 0.78 at the 80% level	**Yes** for the combined score on point estimates. Not for the imaging model at LUMIERE (0.708) or the clinical-history model at UCSF (0.696)	4.3
Flagged scans ranked less well	AUC unflagged minus flagged, CI above 0	Disagreement† **MU only**. Set holding both outcomes **MU and LUMIERE**[a]	4.3
Setting flagged scans aside beats matched removal	Retained minus AUC after matched removal, CI above 0	Disagreement† **MU only**. Set holding both outcomes **UCSF only**[a]	4.3
Reliability card, share of scans in tested situations	None†	Four tested situations, none at UCSF	4.3
Reports			
Reports pass the checker	$\geq 95\%$ in each cohort	First evaluation **No** (every LUMIERE report failed). After one revision **Yes**	4.4
Thin evidence stated	100%	**Yes**	4.4
Reports follow a changed record	≥ 27 of 30 perturbed records	**Yes**	4.4
Checker rejects introduced errors (synthetic records)	All errors of the kinds its rules target	**No**. A wrong tier word was missed	4.4
Usefulness to clinicians	Not tested	A reader study is designed (Supplementary Section S21.1)	5

The plans committed before each test were detection (2b4c927), UCSF-ALPTDG (0e7c9fe), the system (e7cba80), second report evaluation (4f7e7b9), checker mutation test (4f854cb). Matched removal keeps the mix of care phases in the removed scans (random removal at LUMIERE, which has no phase). Over all scans of gliomas of all grades, a planned secondary population, the planned primary model gained +0.041 (+0.006 to +0.081) over the clinical-history model, which may reflect its schedule inputs (Section 4.5).

[a] With each cohort's own medians, the rule the reports use. With medians fixed on MU, disagreement also met both tests at LUMIERE.

4. Validation

4.1. The volumes mostly agreed with a second segmentation but changed between close scans

The released MU volumes mostly agreed with a second automatic segmentation but changed considerably between scans taken close together. A second segmentation of all 594 MU scans (BraTS 2021 SwinUNETR) gave a median Dice coefficient (overlap from 0 to 1) of 0.80 for enhancing tumour, with a quarter of scans below 0.49, and 0.85 for the whole lesion. The log volumes correlated at 0.92 and 0.91. A ridge model on the volumes and their change detected progression no better with the new masks (difference −0.02, −0.06 to +0.03), so the planned test for a gain was not met. Between consecutive scans taken within 30 days of each other, the released enhancing volume changed by a median factor of 2.6, partly through a second resection, and one consistent segmenter reduced this only to 2.2. The report therefore calls the volumes indicative. No expert read the masks again, and the phase of care was taken from the recorded dates without a further check.

4.2. The scores ranked scans on MU and, without refitting, in two external cohorts

The three models ranked scans on MU and kept ranking them, locked, in both external cohorts. In the intended use on MU (146 scans), the clinical-history model detected the progression scan with an AUC of 0.82 (95% CI 0.75–0.90) and the imaging model with 0.71 (0.63–0.78)†. Clinical data alone scored 0.53 (0.45–0.61; Figure 2). The combined score also reached 0.82 (0.74–0.89)†, with no detectable gain over the clinical-history model (−0.001, −0.05 to +0.05). The planned primary model, which added image measurements and three schedule inputs to the clinical-history model, fell short of its target gain of 0.03 (Table 4). Within a phase of care, adding treatment context to clinical data and images gained nothing detectable† (Section 4.5). When every scan before progression counted as a negative, including scans shortly before it, the clinical-history model scored 0.81 (0.76–0.86). Forecasting progression within 120 days was harder, at 0.60 (0.52–0.69) in the intended use (Supplementary Section S23).

The planned reduced form of the combined model, which scored 0.80 (0.73–0.89) on MU by cross validation, reached 0.77 (0.68–0.86) at LUMIERE and met its planned transfer rule (Table 4). Its gain over the reduced clinical-history model† was 0.08 (−0.01 to +0.16). At UCSF both planned imaging models transferred. Images from one scan reached 0.80 (0.69–0.89) and the imaging model 0.73 (0.60–0.84). Locked, the imaging model scored 0.07 below images from one scan (−0.13 to −0.01). Refitted on UCSF, it gained 0.14 (0.03–0.26) over its locked version. The reduced clinical-history model scored 0.56 (0.42–0.69) at UCSF, where its radiotherapy flag is approximated. The locked combined model reached 0.86 (0.76–0.93). Part of this reflects scan timing, because the interval between the two scans was shorter before progression†.

As a probability, the combined score showed no detectable miscalibration in the large at LUMIERE† (observed over expected 0.97, 0.81–1.14) but overpredicted progression at UCSF (0.72, 0.55–0.91), and a threshold at the MU median† called 90% of UCSF scans positive.

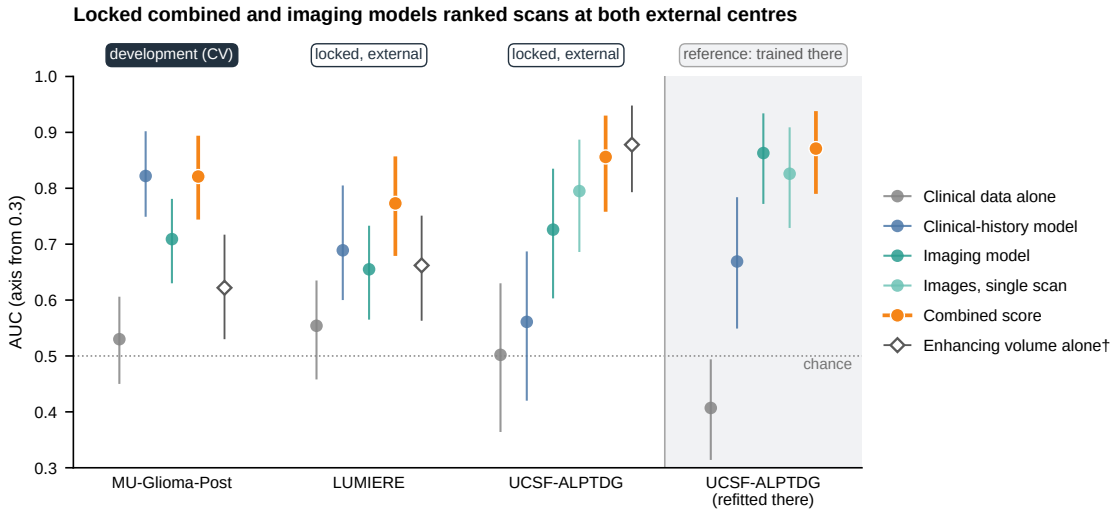


Figure 2. The combined and imaging models rank scans at two external centres without refitting. AUC for detecting the scan at which progression was documented. MU used repeated cross validation in the intended use. LUMIERE and UCSF used models fitted once on MU and applied unchanged, with the reduced clinical-history inputs. The UCSF refit is a reference, where clinical data alone fall below 0.5, consistent with the downward bias of AUC estimated by cross validation without signal. At MU the combined score is its full form. Externally, it has a reduced form. At UCSF the images are shown for images from one scan and for the imaging model. The imaging model and combined score on MU and the clinical data alone, reduced clinical-history and imaging models at LUMIERE were not fixed in advance†. The open diamond is log enhancing tumour volume alone† (at LUMIERE, the 163 scans with enhancing tumour). Points have 95% CIs from bootstrap samples of patients on an axis from 0.3. The dotted line marks chance.

4.3. The combined score's conformal sets kept their coverage on point estimates, and disagreement marked weaker ranking at MU

The combined score's conformal sets met their coverage target on point estimates in every cohort, but coverage diverged by outcome at the external cohorts. The 80% conformal sets of the combined score met the planned coverage threshold on point estimates in all three cohorts, covering 78% to 87% of scans (Figure 3A). The intervals at MU (76% to 88%) and UCSF (68% to 87%) reached below the threshold, and the imaging model's set at LUMIERE and the clinical-history model's set at UCSF fell below it. Coverage differed by outcome at the external centres†. LUMIERE covered 77% (67% to 87%) of progression scans and 94% of scans without progression, and UCSF 97% of progression and 65% (51% to 79%) of scans without progression. Calibrating the sets separately for each outcome on MU did not fix this, and coverage for scans without progression fell to 73% (63% to 82%) even on MU.

Disagreement between the clinical-history and imaging models† marked scans that the combined score ranked less well at MU, but not at the external centres. On MU the combined score's AUC was 0.31 higher where the two models agreed (0.18–0.43), and setting the disagreeing scans aside gained 0.08 (0.05–0.12) over removing as many scans matched for phase of care, meeting both planned tests (Table 4; Figure 3B). With each cohort's own medians, the rule the reports use, neither test was met at LUMIERE or UCSF. With medians fixed on MU, both were met at LUMIERE but not at UCSF (Figure 3B). A conformal set holding both outcomes, used as a flag, met one of the two planned tests in each cohort (Table 4).

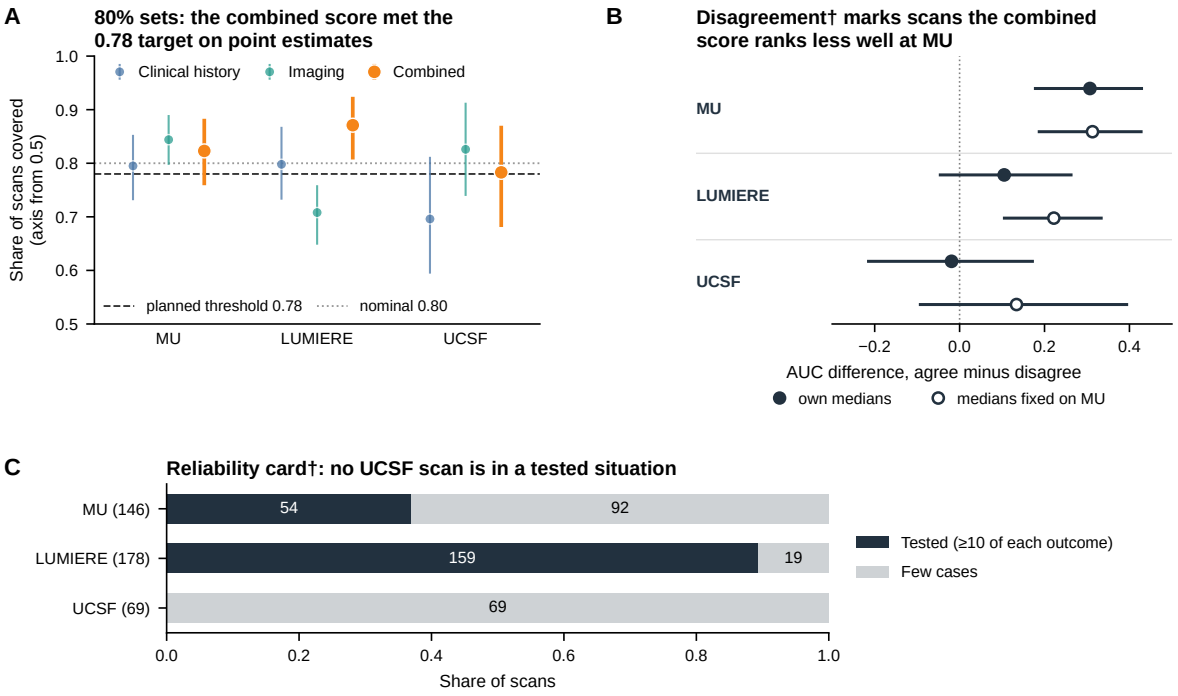


Figure 3. The combined score's conformal sets meet their coverage target on point estimates. Disagreement† marks scans where the combined score ranks less well at MU. (A) Coverage of the 80% conformal sets by cohort and model. The dashed line marks the planned threshold (0.78), and the dotted line marks the nominal level (0.80). (B) The combined score's AUC where the two models agree minus where they disagree, with each cohort's own medians (filled) and with medians fixed on MU (open, using the other folds' medians on MU). The tests were planned, but the median rules were not†. (C) Scans whose situation reached the reliability card's† tested tier. The rest had few cases, and none was untested.

Only four situations on the reliability card† reached the tested tier (Figure 3C), one at MU and three at LUMIERE. LUMIERE records only the rater's flag for the first three months, so its situations differ by enhancement only. No scan in these sets was in an untested situation. Every UCSF situation had few cases, and at MU no situation in the first 12 weeks after radiotherapy held more than three progression scans. In a descriptive comparison, tested situations had lower error than situations with few cases at LUMIERE (29% against 42% of scans misclassified at the cohort median) but not at MU (29% against 24%), where the scans with few cases are mostly clearly without progression early after radiotherapy.

For a quarter to over a third of scans, the confidence statements flag something a score alone does not show (Table 5), and the flagged scans were misclassified more often than the rest.

Table 5. What a score alone does not show. Share of scans in each cohort for which each confidence statement flags a doubt, and the combined score's error when every scan above the cohort's median score is called progression. † not fixed in advance.

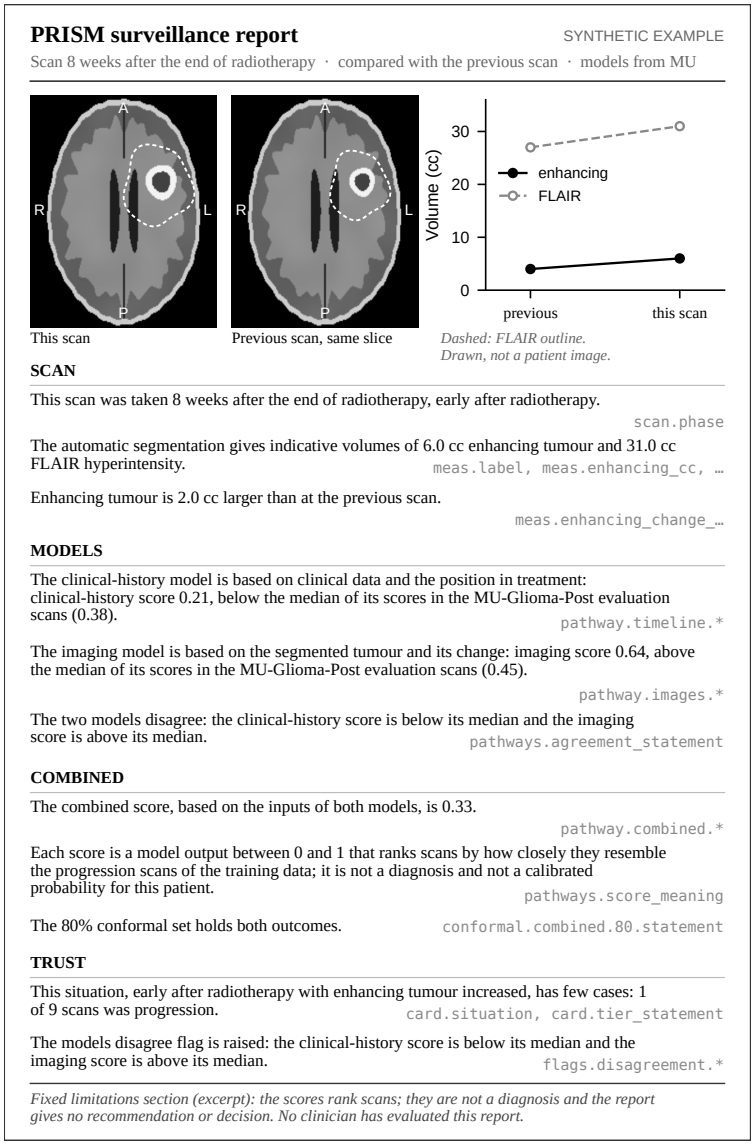
	MU	LUMIERE	UCSF
Conformal set holds both outcomes	25%	38%	36%
Clinical-history and imaging models disagree†	40%	57%	55%
Situation with few cases on the reliability card†	63%	11%	100%
Error where the set holds both outcomes†	37%	48%	32%
Error on the other scans†	22%	20%	25%

4.4. After one revision of the prompt, every report passed the checker

Every report written by Qwen3.8-27B for the scans of patients not used to revise the prompt passed the checker, 114 at MU, 162 at LUMIERE and 130 at UCSF. We gave six other local models the same prompt, revised for Qwen3.8-27B, and only Gemma-4 31B also had every report pass (Figure 4B). All 220 reports in situations with few cases or untested situations said so (71 at MU, 19 at LUMIERE and 130 at UCSF). When one confidence or measurement field was changed in 30 records from the three cohorts, 28 reports followed the change, against a planned minimum of 27 (Table 4). Before the revision of the prompt†, every LUMIERE report failed because a fixed sentence quoted a number held in no field. We corrected the prompt, left the checker unchanged and evaluated again on patients not used for the revision.

The checker caught 1,252 of 1,280 targeted errors in templated synthetic reports. We built correct reports from the unit test templates for 52 synthetic records and introduced one error at a time of the kinds listed in Section 2.4. The checker rejected none of the correct reports and rejected most reports with an error. All 28 misses wrote "tested" into the tier sentence of a situation with few cases. Because no rule covers them, the checker also accepted every report in which a direction word was flipped or the enhancing and FLAIR volumes were swapped between their labels. The checker cannot tell whether the segmentation itself is right.

A **A PRISM report (synthetic example)**



B **Reports passing the check**

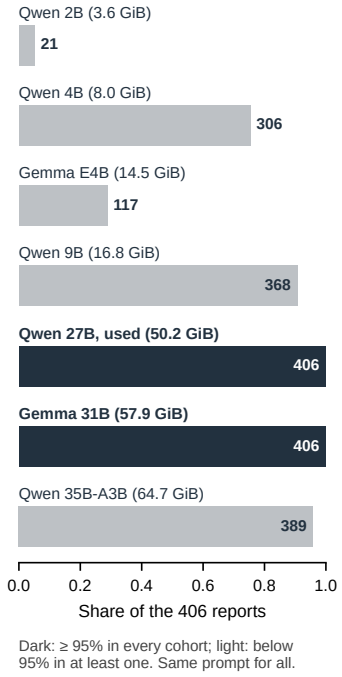


Figure 4. What the clinician receives, and the check. (A) A report for a synthetic scan eight weeks after radiotherapy with growing enhancement (Section 5). The segmentation figure is drawn here rather than taken from a patient, and the narrative is synthetic. Volumes and scores are invented, and the card row is MU's. Each sentence lists, in grey, the record fields it cites. The delivered report also holds tables of every value and a fixed limitations section (excerpt shown). (B) Reports passing the check, of 406, for seven local language models given the same prompt (Qwen3.5 2B, 4B and 9B, Gemma-4 E4B and 31B, Qwen3.8-27B used in PRISM, and Qwen3.5-35B-A3B), after one revision of the prompt† (Section 4.4). Dark bars show models for which at least 95% passed in every cohort, and the model weight size is in brackets.

4.5. The scores rested on the treatment schedule at MU and on enhancement at UCSF

At MU the clinical-history and combined scores depended mainly on where the patient was in treatment, and at UCSF the imaging scores followed the volume of enhancing tissue† (Figure 5). At MU most progression scans came more than 12 weeks after radiotherapy and most scans without progression within those 12 weeks (Table 2), so the schedule alone could separate most of them. An AUC within a phase compares only scans from the same phase, and most such comparisons come from the later phase (56 progression and 32 scans without progression). Adding treatment context to clinical data and images† changed the AUC by +0.11 (0.04–0.20) across phases of care and by −0.02 (−0.13 to +0.11) within them (Figure 5A). Within a phase, the imaging model scored 0.68 (0.58–0.77) and the clinical-history model 0.64 (0.51–0.78), with no difference detectable. The interval within a phase is wide enough that a gain as large as the one across phases cannot be ruled out.

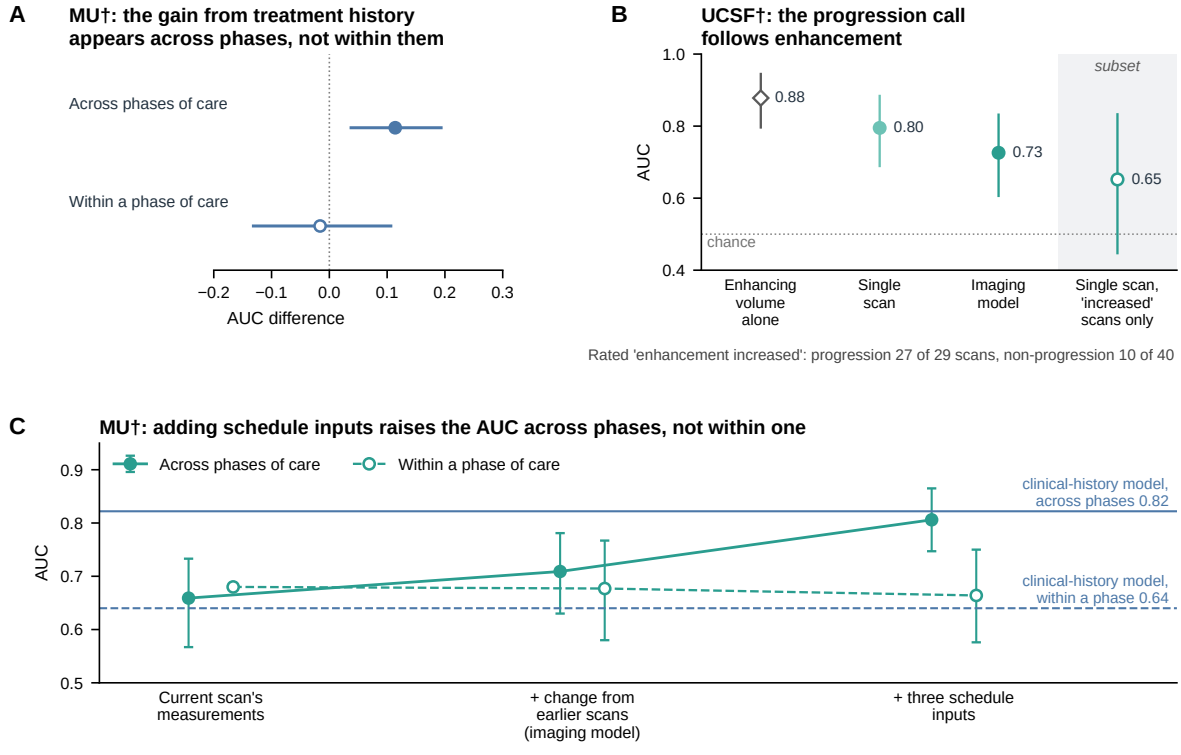


Figure 5. The scores followed the schedule at MU and enhancement at UCSF. (A) At MU, the gain from adding treatment context to clinical data and images, across and within phases of care. (B) At UCSF, enhancing volume alone beside the locked imaging models. The open point is restricted to the scans rated as showing increased enhancement. (C) At MU, the imaging model built in three steps, across and within phases of care, against the clinical-history model. The third step adds the planned primary model's schedule inputs to the image inputs, without its clinical-history inputs. The gain at the second step has an interval that includes zero, and the first point within a phase has no interval. The audits were not fixed in advance (†), but the locked imaging models' AUCs over all 69 scans in (B) were planned. Points with 95% CIs from bootstrap samples of patients.

On MU the imaging model gained most from inputs that carry timing†. Adding change from earlier scans to the current scan's measurements, which gives the imaging model, raised the AUC by 0.05, with an interval that included zero (−0.003 to +0.106). Adding the planned primary model's three schedule inputs (Table 1) raised it by a further 0.10 (0.05–0.15) across phases of care but not

within them (−0.01, −0.06 to +0.04) (Figure 5C).

At UCSF the progression labels coincided with a rated increase in enhancement†, so a model could score well there by tracking it. UCSF rated enhancement as increased at 27 of the 29 progression scans, and progression was dated a median of 7 days after the scan (Figure 5B). Enhancing volume alone reached 0.88 (0.79–0.95) there. This was above the point estimates of the locked imaging models and above its own AUC at MU and LUMIERE (Figure 2). Among the 37 scans rated as increased, images from one scan reached 0.65 (0.44–0.84), so this sample cannot show whether the measurements add to that rating. Deep networks tuned on UCSF scans beat the control with permuted labels (Table 4), and their scores tracked enhancing volume†.

5. Discussion

PRISM's combined and imaging models ranked scans at two other centres without refitting, but the information a score used depended on how each cohort was scanned and labelled. At MU, the score mostly followed the treatment schedule and at UCSF it followed the enhancement on which progression was called†. An AUC cannot show which evidence a model used. The audits of Section 4.5 suggested which evidence each model used, and the report shows the clinical-history and imaging scores side by side, so a disagreement between them is visible at each scan. Process variables can carry much of the apparent skill of an image model (Badgeley et al. 2019), and image models for glioblastoma surveillance have been benchmarked separately at the first and second scans after radiotherapy (Guo and Mirzaei 2025). Outside oncology, clinical data were the strongest input in one Alzheimer's cohort and images in another (Brima et al. 2026), and models with equal AUC can behave differently on new data (D'Amour et al. 2022). We found no earlier study that compared image and treatment context evidence within phases of care in glioma surveillance, or showed that what carries a locked score differs between externally tested cohorts. The risk scale did not transfer to UCSF, so a new centre would recalibrate the score and set any threshold on local scans.

Because its modules are separate, PRISM runs on whatever a cohort holds, provided it has scans and masks. The imaging model needs only those, which most public datasets have, and images from one scan serve where there is no earlier scan. The clinical-history model is added where dates exist. The reliability card is a lookup table, so a new cohort adds rows to it and a situation's tier can change as evidence accumulates.

In the weeks when pseudoprogression is most common, PRISM has too few comparable scans to support a call, and its report says so. Consider a scan eight weeks after radiotherapy in which the enhancing tumour has grown (Figure 4A). The clinical-history model scores such a scan low, because most scans at this point in treatment are not progression, and the imaging model scores it high, because enhancement has grown. The report shows both, states that they disagree and that the conformal set holds both outcomes, and records from the card that the situation has few cases (at MU, 1 of 9 such scans was progression†). The clinician then decides with information PRISM does not use, such as symptoms, steroid dose and the radiotherapy plan.

Each component of PRISM has been described before, from longitudinal measurement of change (Rudie et al. 2022) and late fusion of image and clinical models (Boehm et al. 2022; Gomaa et al. 2025) to conformal prediction in neuro oncology imaging (Elyassirad et al. 2025; Angelopoulos et al. 2024), disagreement between modalities (Mahapatra et al. 2026), reliability from similar cases (Habas et al. 2007; Mitchell et al. 2019) and reports written by local language models from

tumour measurements (Heras Rivera et al. 2026; Valerio et al. 2025). The nearest systems combine some of these. One assigns a BT-RADS category at follow up from tumour volumes and treatment variables extracted by a language model, at one centre and without an uncertainty measure for each scan (Jabal et al. 2026). Another keeps a ledger of evidence with a confidence for each claim in longitudinal neuro oncology, but works from reports rather than images and checks its evidence with a language model (Liu et al. 2026). On these cohorts, earlier work predicted a label for each patient on MU (Christodoulou et al. 2026) and response categories from longitudinal MRI (Tikhonov et al. 2025), and a model trained on MU was applied unchanged at UCSF to predict survival (Ding et al. 2026). Chen et al. (2026) predicted recurrence risk from preoperative and postoperative MRI across centres. None reports image and treatment context evidence separately at each scan, or tests a progression call for each scan locked at two external centres. For reports, an ECG system already combines a local language model, a verifier and a conformal set (Pham et al. 2026), but it forbids numbers in the text rather than checking each against a typed record. PRISM therefore combines known parts. It shows image and treatment context evidence at each scan with a confidence from measured performance, and we tested the combination at two external centres under rules fixed in advance.

PRISM was developed on one public cohort, and the two external tests are its only evidence that the ranking transfers. The label is progression as each centre documented it. No scan was read again, and at UCSF most progression scans were also rated as showing increased enhancement. Of the confidence measures, only the combined score's average coverage met its planned threshold at all three centres, and only on point estimates. Disagreement† marked scans the combined score ranked less well at MU alone. That rule needs a batch of a cohort's scans, and its single scan form (MU medians) was defined after the results had been seen. No UCSF scan fell in a tested situation. The conformal set comes from a refitted copy of the combined model, and the card counts the scan it describes. Calibrating the displayed model itself and building the card without the scan are the next fixes, and both need the external cohorts' scores again. The checker misses a wrong tier word and numbers swapped between labels, and no person has checked a sample of the real reports for such errors. A fixed template could state the same fields, and we have not tested whether the language model's wording helps a reader more than a template does. Subgroups are small, the analyses marked † need testing on new data, and no clinician has yet read the reports. PRISM forecast progression within 120 days less well than it detected progression at the scan (AUC 0.60 in the intended use).

The next steps are to pool the rare situations (early progression after radiotherapy, and growing enhancement without progression) across cohorts, including the Burdenko and UCSD-PTGBM datasets, and to test models tuned on them there. A reader study would show whether separate evidence and a stated confidence change clinicians' decisions.

Conclusion

PRISM keeps evidence from images and treatment context apart and shows the clinician both scores at every scan with how far each has been tested there. Its combined and imaging models ranked scans at two other centres without refitting. Audits showed that the scores followed the treatment schedule at MU and enhancement at UCSF. An AUC alone does not reveal this difference.

Acknowledgements

[PENDING AUTHOR DECISION].

Disclosure statement

[PENDING AUTHOR DECISION: competing interests, or "The authors report there are no competing interests to declare."] If any author serves on the editorial or advisory board of this journal, state it here.

Declaration of generative AI use

[DRAFT, PENDING AUTHOR CONFIRMATION:] Generative AI is part of the system studied. PRISM's report generator uses a local language model (Qwen3.8-27B; six other models with open weights in a comparison); all ran offline on the authors' computing cluster and saw only derived system records. An AI assistant (Claude, Anthropic) was used to write and run analysis code and to draft and edit the manuscript under the authors' direction; it was given no data at the patient level through any hosted service. The authors reviewed all code, results and text and take full responsibility for the content of this article.

Funding

No funding was obtained for this work.

Ethics statement

[PENDING AUTHOR DECISION: secondary analysis of datasets with identifiers removed (MU-Glioma-Post and LUMIERE, public; UCSF-ALPTDG under a data use agreement); state the institutional determination or waiver and cite the source datasets' approvals.] No registered protocol. Analysis plans were committed to the repository before each external test; their commits are listed in Table 4. Patient and public involvement: none.

Data availability statement

[PENDING AUTHOR DECISION: MU-Glioma-Post (TCIA, doi:10.7937/7K9K-3C83) and LUMIERE (figshare; Suter et al. 2022) are public; UCSF-ALPTDG is available from its provider under a data use agreement; state which derived tables, fold assignments and locked model files are released, with a DOI.] Code: [PENDING: public repository link and release tag.]

Supplemental material

Supplemental data for this article are provided as a separate file.

Notes on contributors

[PENDING: a short biographical note for each author (no more than 200 words each).]

References

An C et al. 2021. Radiomics machine learning study with a small sample size: Single random training-test set split may lead to unreliable results. PLOS ONE. 16(8):e0256152. https://doi.org/10.1371/journal.pone.0256152

Angelopoulos AN, Bates S. 2023. Conformal prediction: A gentle introduction. Found Trends Mach Learn. 16(4):494–591. https://doi.org/10.1561/2200000101

Angelopoulos AN et al. 2024. Conformal triage for medical imaging AI deployment. [Preprint, medRxiv]

Badgeley MA, Zech JR, Oakden-Rayner L et al. 2019. Deep learning predicts hip fracture using confounding patient and healthcare variables. npj Digit Med. 2:31. https://doi.org/10.1038/s41746-019-0105-1

Boehm KM, Aherne EA, Ellenson L, Nikolovski I et al. 2022. Multimodal data integration using machine learning improves risk stratification of high-grade serous ovarian cancer. Nat Cancer. 3:723–733. https://doi.org/10.1038/s43018-022-00388-9

Brandes AA, Franceschi E, Tosoni A et al. 2008. MGMT promoter methylation status can predict the incidence and outcome of pseudoprogression after concomitant radiochemotherapy in newly diagnosed glioblastoma patients. J Clin Oncol. 26(13):2192–2197. https://doi.org/10.1200/JCO.2007.14.8163

Brima Y, Atemkeng M, Namamula LR, Vacavant A. 2026. A multimodal explainable deep learning framework for Alzheimer's disease diagnosis using 3D magnetic resonance imaging and clinical data. [Arxiv:2609.12410].

Chen C, Guo F, Huang D et al. 2026. Predicting recurrence risk of glioblastoma based on preoperative-postoperative longitudinal MRI: A multicenter study. Bioengineering. 13(6):668. https://doi.org/10.3390/bioengineering13060668

Christodoulou RC, Vamvouras G, Pitsillos R, Solomou EE, Georgiou MF. 2026. Explainable radiomics with probability calibration for postoperative glioblastoma surveillance. Eur J Radiol Artif Intell. 5:100074. https://doi.org/10.1016/j.ejrai.2026.100074

D'Amour A et al. 2022. Underspecification presents challenges for credibility in modern machine learning. J Mach Learn Res. 23(226):1–61.

Ding T, Zou Y, Chen C, Shah M, Tian Y. 2026. CLARITY: Medical world model for guiding treatment decisions by simulating context-aware disease trajectories. [Preprint, arXiv:2512.08029]. `https://arxiv.org/abs/2512.08029`

Elyassirad D, Gheiji B, Vatanparast M, Ahmadzadeh AM, Faghani S. 2025. CONSeg: Voxelwise uncertainty quantification for glioma segmentation using conformal prediction. Am J Neuroradiol. 46(12):2553–2560. https://doi.org/10.3174/ajnr.A8914

Fields BKK et al. 2024. The University of California San Francisco adult longitudinal post-treatment diffuse glioma MRI dataset. Radiol Artif Intell. 6(4):e230182. https://doi.org/10.1148/ryai.230182

Gomaa A, Huang Y, Stephan P et al. 2025. A self-supervised multimodal deep learning approach to differentiate post-radiotherapy progression from pseudoprogression in glioblastoma. Sci Reports. 15:17133. https://doi.org/10.1038/s41598-025-02026-7

Gore A et al. 2019. Institutional implementation of a structured reporting system: Our experience with the Brain Tumor Reporting and Data System. Acad Radiol. 26(7):974–980. https://doi.org/10.1016/j.acra.2018.12.023

Guo W, Mirzaei G. 2025. Timepoint-specific benchmarking of deep learning models for glioblastoma follow-up MRI. Cancers. 18(1):36. https://doi.org/10.3390/cancers18010036

Habas PA, Zurada JM, Elmaghraby AS, Tourassi GD. 2007. Reliability analysis framework for computer-assisted medical decision systems. Med Phys. 34(2):763–772. https://doi.org/10.1118/1.2432409

Heras Rivera JE et al. 2026. BTReport: A framework for brain tumor radiology report generation with clinically relevant features. [Preprint, arXiv:2602.16006]. https://arxiv.org/abs/2602.16006

Jabal MS et al. 2026. Agentic automation of BT-RADS scoring: End-to-end multi-agent system for standardized brain tumor follow-up assessment. [Preprint, arXiv:2603.21494]. https://arxiv.org/abs/2603.21494

Kwon W et al. 2023. Efficient memory management for large language model serving with PagedAttention. In: Proceedings of the 29th Symposium on Operating Systems Principles. ACM. p 611–626

Liu Y et al. 2026. TumorBoard: Evidence-grounded multi-agent decision support for longitudinal neuro-oncology. [Preprint, arXiv:2608.03190]. https://arxiv.org/abs/2608.03190

Mahapatra D et al. 2026. EnTrust: Modeling inter-modal conflict for trustworthy multimodal medical image analysis. [Preprint, arXiv:2606.21384]. https://arxiv.org/abs/2606.21384

Mahmoud E, Gass J, Dhemesh Y et al. 2025. MU-Glioma Post: A comprehensive dataset of automated MR multi-sequence segmentation and clinical features. Sci Data. 12(1):1847. https://doi.org/10.1038/s41597-025-06011-7

Mitchell M, Wu S, Zaldivar A et al. 2019. Model cards for model reporting. In: Proceedings of the Conference on Fairness, Accountability, and Transparency (FAT* '19). ACM. p 220–229

Pedregosa F et al. 2011. Scikit-learn: Machine learning in Python. J Mach Learn Res. 12:2825–2830.

Pham KD et al. 2026. Evidence-augmented conformal risk control with verifier-gated LLM reporting for trustworthy ECG decision support. In: Proceedings of the IEEE Engineering in Medicine and Biology Conference (EMBC). Author copy at https://cs.emory.edu/ jyang71/files/ea-crc.pdf.

Romano Y, Barber RF, Sabatti C, Candès E. 2020. With malice toward none: Assessing uncertainty via equalized coverage. Harv Data Sci Rev. 2(2). https://doi.org/10.1162/99608f92.03f00592

Rudie JD, Calabrese E, Saluja R et al. 2022. Longitudinal assessment of posttreatment diffuse glioma tissue volumes with three-dimensional convolutional neural networks. Radiol Artif Intell. 4(5):e210243. https://doi.org/10.1148/ryai.210243

Suter Y, Knecht U, Valenzuela W et al. 2022. The LUMIERE dataset: Longitudinal glioblastoma MRI with expert RANO evaluation. Sci Data. 9(1):768. https://doi.org/10.1038/s41597-022-01881-7

Taal W, Brandsma D, de Bruin HG et al. 2008. Incidence of early pseudo-progression in a cohort of malignant glioma patients treated with chemoirradiation with temozolomide. Cancer. 113(2):405–410. https://doi.org/10.1002/cncr.23562

Tikhonov D et al. 2025. Predicting brain tumor response to therapy using a hybrid deep learning and radiomics approach. [Arxiv:2509.06511]. Submitted to the BraTS-Lighthouse 2025 Challenge.

Valerio AG, Trufanova K, de Benedictis S, Vessio G, Castellano G. 2025. From segmentation to explanation: Generating textual reports from MRI with LLMs. Comput Methods Programs Biomed. 270:108922. https://doi.org/10.1016/j.cmpb.2025.108922

Varoquaux G. 2018. Cross-validation failure: Small sample sizes lead to large error bars. NeuroImage. 180:68–77. https://doi.org/10.1016/j.neuroimage.2017.06.061

Weinberg BD et al. 2018. Management-based structured reporting of posttreatment glioma response with the Brain Tumor Reporting and Data System. J Am Coll Radiol. 15(5):767–771. https://doi.org/10.1016/j.jacr.2018.01.022

Wen PY, van den Bent M, Youssef G et al. 2023. RANO 2.0: Update to the Response Assessment in

Neuro-Oncology criteria for high- and low-grade gliomas in adults. J Clin Oncol. 41(33):5187–5199. https://doi.org/10.1200/JCO.23.01059

Zeba M, Al Mamun A, Tithee KJ, Sutradhar D et al. 2026. Mitigating hallucinations in healthcare LLMs with granular fact-checking and domain-specific adaptation. Expert Syst Appl. 329. Preprint arXiv:2512.16189.

Zhang JY et al. 2020. Quantitative improvement in brain tumor MRI through structured reporting (BT-RADS). Acad Radiol. 27(6):780–784. https://doi.org/10.1016/j.acra.2019.07.028